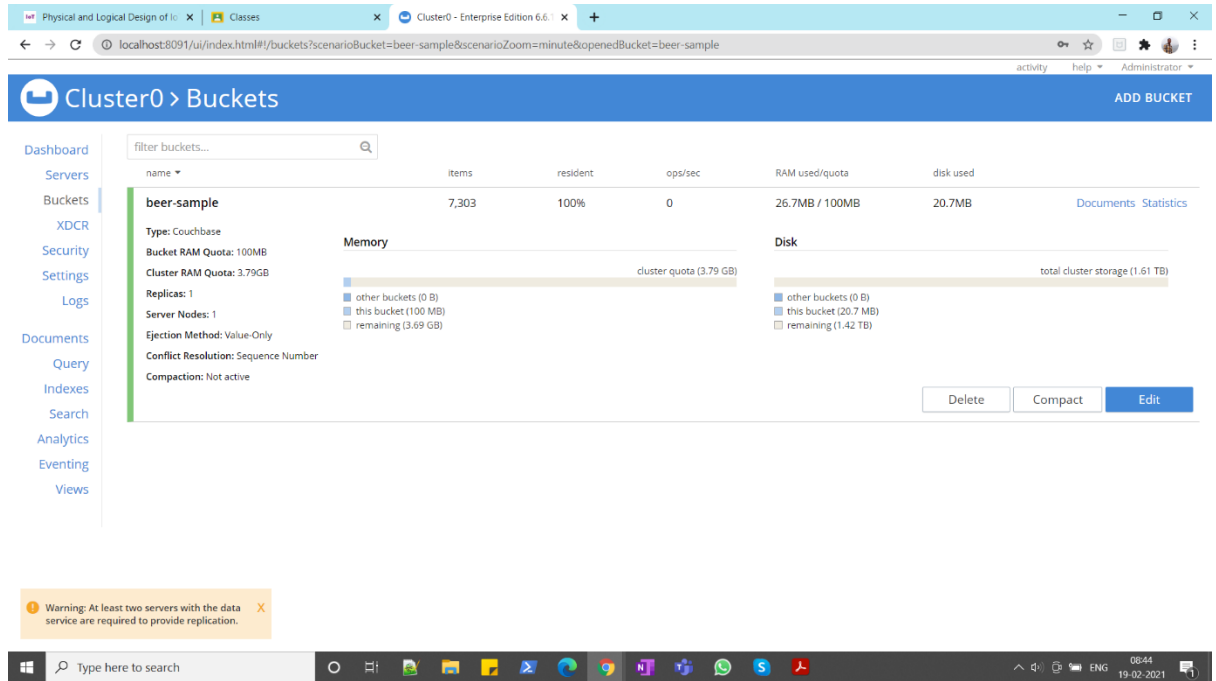


NoSQL Assignment:

1.) Import 'beer-sample' Bucket.



Cluster0 > Buckets

filter buckets...

name	items	resident	ops/sec	RAM used/quota	disk used
beer-sample	7,303	100%	0	26.7MB / 100MB	20.7MB

Type: Couchbase
Bucket RAM Quota: 100MB
Cluster RAM Quota: 3.79GB
Replicas: 1
Server Nodes: 1
Ejection Method: Value-Only
Conflict Resolution: Sequence Number
Compaction: Not active

Memory

- other buckets (0 B)
- this bucket (100 MB)
- remaining (3.69 GB)

Disk

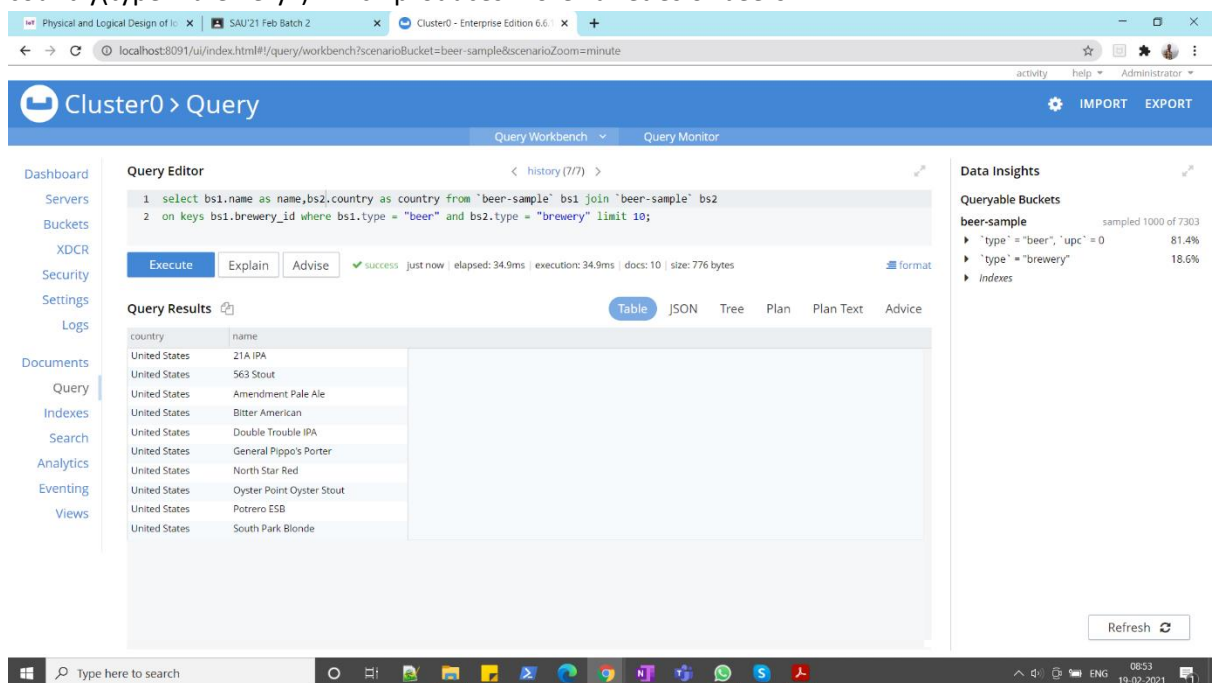
- other buckets (0 B)
- this bucket (20.7 MB)
- remaining (1.42 TB)

total cluster storage (1.61 TB)

Delete Compact Edit

Warning: At least two servers with the data service are required to provide replication.

2.) Write a join query to fetch the Top 10 brewery(type="beer") and their country(type="brewery") which produces more varieties of beers.



Cluster0 > Query

Query Editor

```
1 select bs1.name as name, bs2.country as country from 'beer-sample' bs1 join 'beer-sample' bs2
2 on keys bs1.brewery_id where bs1.type = "beer" and bs2.type = "brewery" limit 10;
```

Execute Explain Advise success just now elapsed: 34.9ms execution: 34.9ms docs: 10 size: 776 bytes

Query Results

country	name
United States	21A IPA
United States	563 Stout
United States	Amendment Pale Ale
United States	Bitter American
United States	Double Trouble IPA
United States	General Pippo's Porter
United States	North Star Red
United States	Oyster Point Oyster Stout
United States	Potrero ESB
United States	South Park Blonde

Data Insights

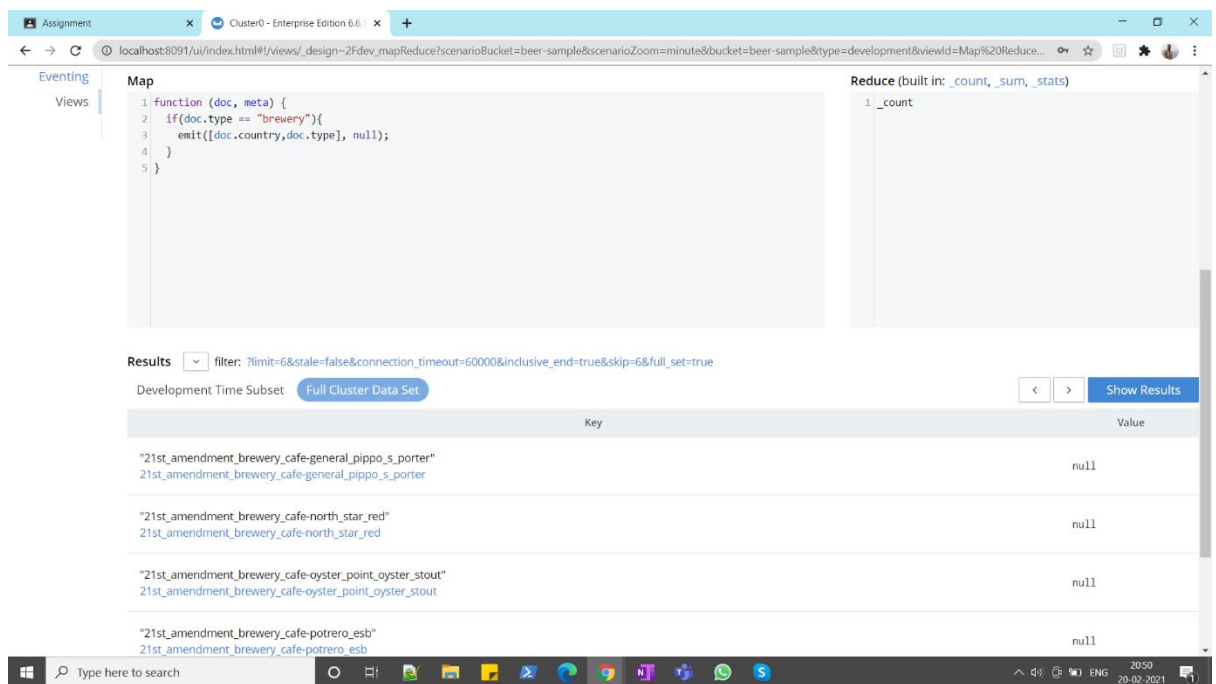
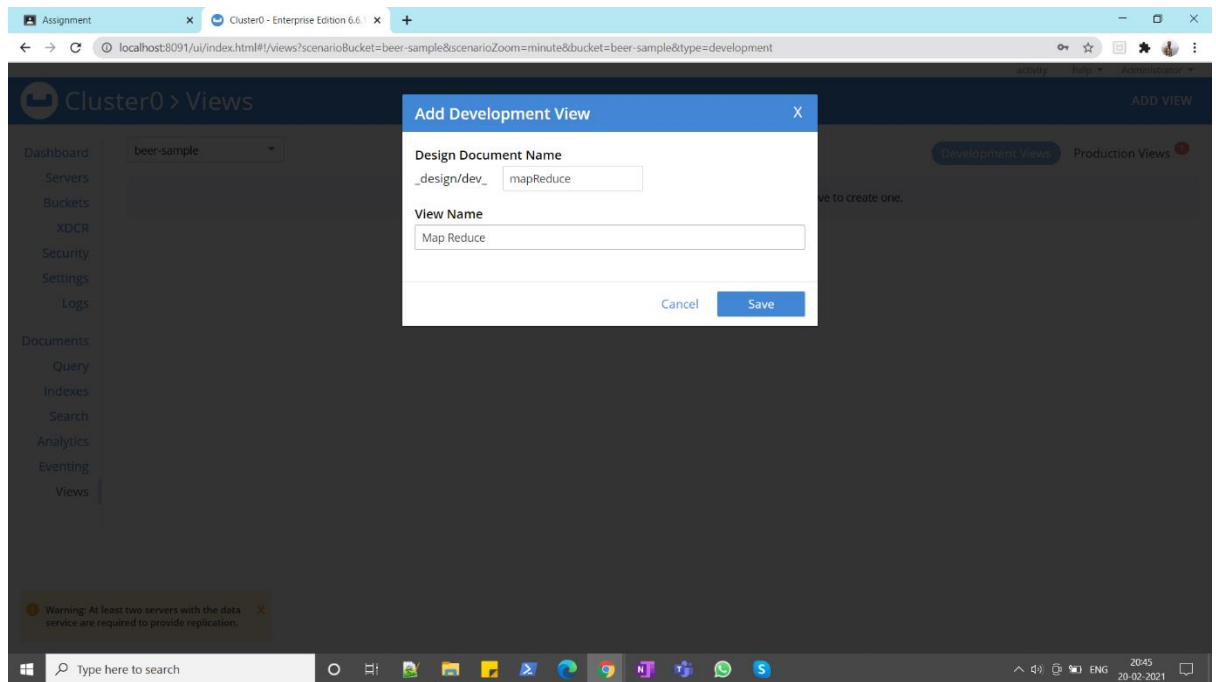
Queryable Buckets

beer-sample sampled 1000 of 7303

- 'type' = "beer", 'upc' = 0 81.4%
- 'type' = "brewery" 18.6%
- indexes

Refresh

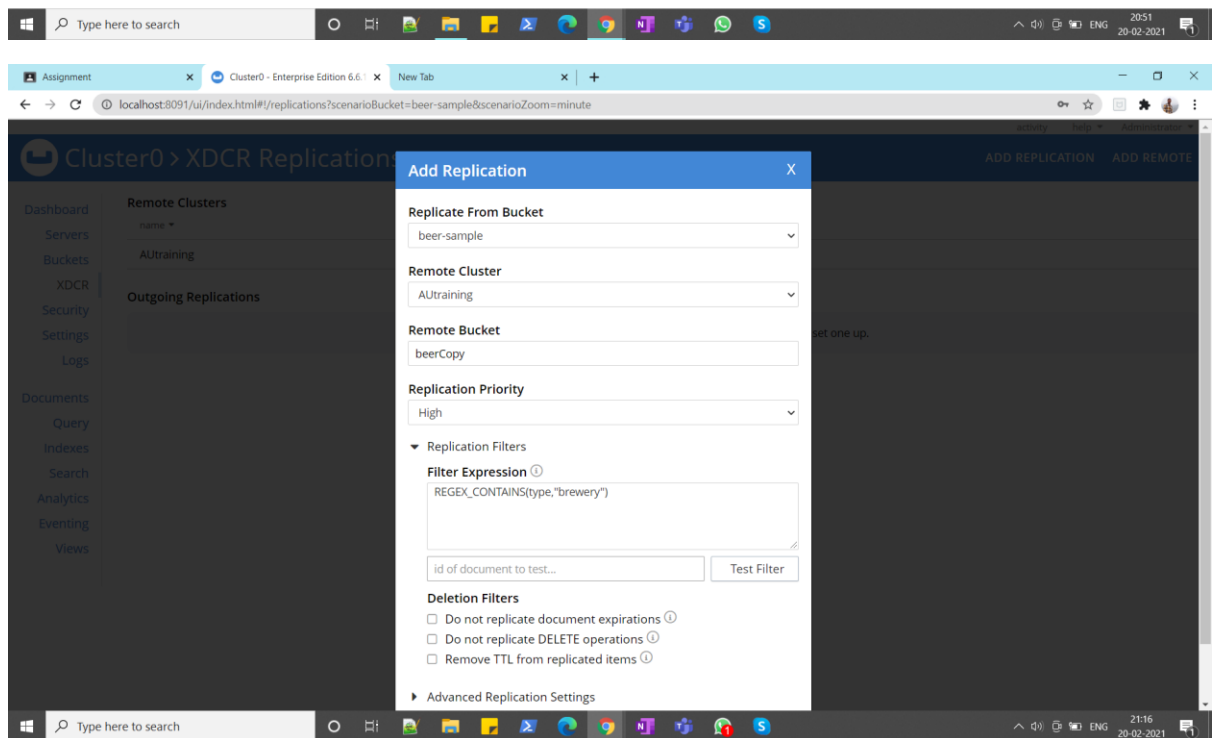
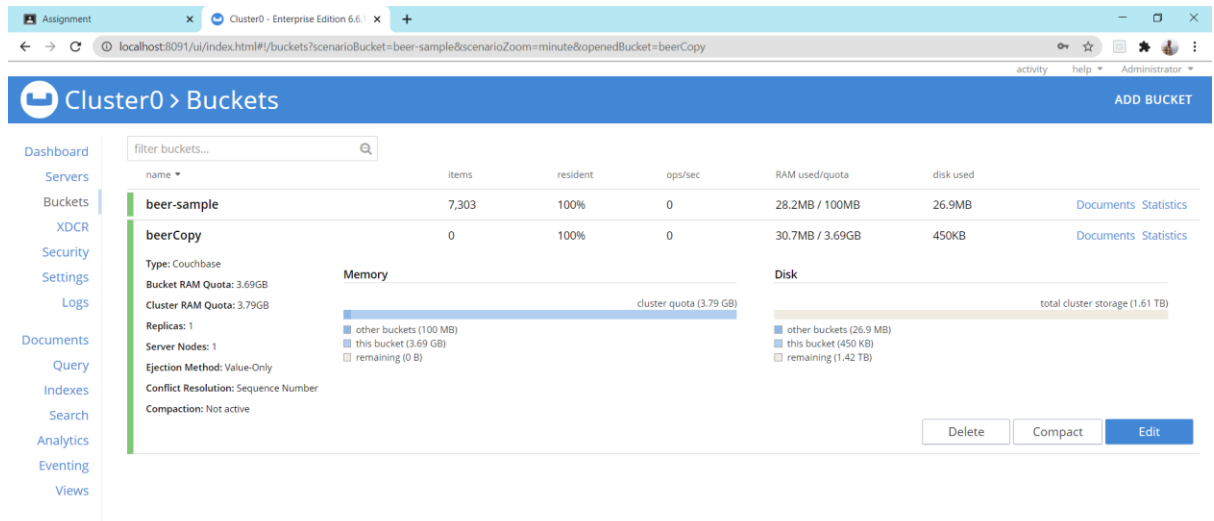
3.) Write a mapreduce to get the number of breweries based on country. Please attach the mapreduce code and json output screenshot.



4.) XDCR:

a) Add a new bucket "Brewery".

b) Create a XDCR with a filter(type='brewery') to replicate only the brewery entity from `beer-sample` bucket.



5.) CLI:

a) Add a new bucket "Beer".

b) Using CLI - do a cbexport of the entire 'beer-sample'

c) And do a cbimport with "brewery_id" as primary key. As a result, in the new bucket - only "beer" documents will be imported with their respective brewery name as meta().id

d) Share your observation on counts with screenshot.

```

--no-ssl-verify      Skips SSL verification of certificates against CA
-l,--log-file        Specifies a file to log debugging information
-v,--verbose         Print debugging information (Has no effect if a log
                     file is specified)
-t,--threads         The amount of parallelism use (Default is 1)
-h,--help            Prints the help message

c:\Program Files\Couchbase\Server\bin>cbexport json -c couchbase://127.0.0.1 -u Administrator -p Baadsha_1 -b beer-sample -o
C:/Users/"MD SHADAAB FARHAN"/Desktop/beer_sample.json" -f lines -t 4
Json exported to `C:/Users/MD SHADAAB FARHAN/Desktop/beer_sample.json` successfully
Documents exported: 7303 Documents skipped: 0

c:\Program Files\Couchbase\Server\bin>cd

```